

# **MapReduce**

## **Simplified Data Processing on Large Clusters**

**Paper by Jeffrey Dan and Sanjay Ghemawat, Google Inc**

# Outline

- The original MapReduce paper<sup>1</sup> (2004) and the journal version <sup>2</sup>
- Influence of MapReduce paper (2004-2015)
- Beyond MapReduce (2015 - ...)

<sup>1</sup> Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: Simplified Data Processing on Large Clusters." To appear in OSDI (2004): 1.

<sup>2</sup> Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: simplified data processing on large clusters." Communications of the ACM 51.1 (2008): 107-113.

***MapReduce is a **programming model** and  
an associated implementation for  
processing and generating large data sets.***

— Literally the first sentence of the paper

# Contributions

1. Programming model
2. MapReduce implementation (description)

*The major contributions of this work are **a simple and powerful interface** that enables automatic parallelization and distribution of large-scale computations, combined with **an implementation** of this interface that achieves high performance on **large clusters of commodity PCs***

# The **restricted** programming model

map:  $(k1, v1) \Rightarrow \text{List}[(k2, v2)]$   
reduce:  $(k2, \text{List}[v2]) \Rightarrow \text{List}[v3]$

Inspired by LISP and other functional programming

# Classical example

## WordCount

Given lines of text

```
function map( name:String, line:String ) = {  
    for( word in line )  
        emit ( word, 1 )  
}
```

```
function reduce( word:String, counts:Iterator[Int] ) = {  
    sum = 0  
    for( count in counts )  
        sum += count  
    emit ( word, sum )  
}
```

# Monoid requirements

- The reduce function must be associative  
 $(x \cdot y) \cdot z = x \cdot (y \cdot z)$
- The reduce function must have a neutral element  $e$   
 $(x \cdot e) = (e \cdot x) = x, \forall x$

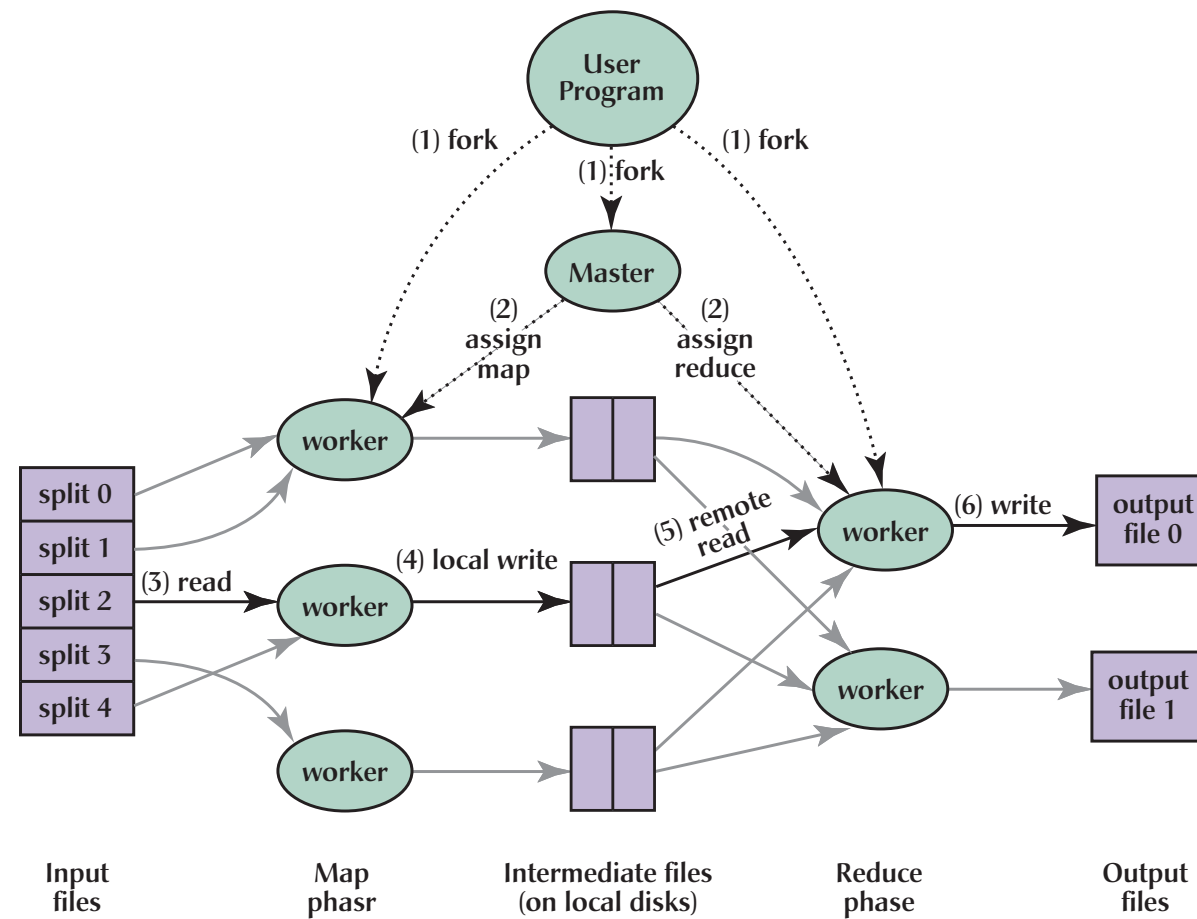


# **It's still a restricted model**

Some operations are not so easy to express in terms of map and reduce.

# The implementation

- Automatic parallelization and distribution
- Fault tolerance
- I/O Scheduling
- Status and monitoring



Taken from <sup>2</sup>

<sup>2</sup> Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: simplified data processing on large clusters." Communications of the ACM 51.1 (2008): 107-113.

# Fault tolerance

- Monitor execution
- Re-execute stale / failed jobs
- Skipping Bad Records


# Optimizations

- Partitioning functions
- Combiner functions

# Influence of MapReduce paper: Birth of an industry

## MapReduce popularity

Relative searches for "MapReduce" on Google



# Yahoo builds Hadoop

- Implements MapReduce framework based on MapReduce paper<sup>3</sup>
- Implements HDFS based on the GFS paper<sup>4</sup>


<sup>3</sup> Hadoop: Open Source implementation of MapReduce

<sup>4</sup> Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung. "The Google file system." ACM SIGOPS operating systems review. Vol. 37. No. 5. ACM, 2003.



## MapReduce popularity

Relative searches for "MapReduce" on Google



# Hadoop ecosystem

PIG, Hive, ZooKeeper and others

*All open source.* Most part of the Apache Software Foundation.

Startups: HortonWorks & Cloudera

# Beyond MapReduce

# MapReduce, the bad parts

- Everything written to disk
  - Slow
  - RAM becomes cheaper (4GB in original paper)
  - Machine Learning workloads
- Very **restricted** programming model
  - Map - Reduce all the things

# Some contenders


- Cascading
  - Wrapper around MapReduce (Google might have something similar)
- Spark
  - Keep the functional bits, throw out the rest

# Spark, the MapReduce killer

- Lazy evaluation
- In memory caching
- More expressive API
- Transition to streaming

## MapReduce popularity

Relative searches for "MapReduce" on Google



# Classical Example revisited

## Wordcount

```
val lines = spark.textFile( "content.txt" )
val counts = lines.flatMap { line =>
    line.split( ' ' ).map( word => ( word, 1 ) )
}.reduceByKey( _ + _ )

println( counts.collect() )
```



# Spark, the MapReduce heir

- Still part of Hadoop ecosystem
- Can run on YARN (MapReduce 3.0)
- Very heavily influenced by the functional programming API of Scala

# Legacy of MapReduce

## Big Data



## Functional Programming